Eshwar Chandrasekharan

Assistant Professor, University of Illinois Urbana-Champaign
eshwar AT illinois DOT edu | eshwarchandrasekharan.com

ACADEMIC APPOINTMENTS

2020 – present **Assistant Professor**, Siebel School of Computing and Data Science, University of Illinois Urbana-Champaign, Urbana, IL

RESEARCH INTERESTS

My research examines how digital environments, from social platforms to AI systems, shape human behavior and access to information. I build empirical methods and human-centered AI systems to make these environments safer and healthier, drawing on causal inference, algorithmic auditing, natural language processing, and system design.

Topics: Human-centered AI, Social Computing, AI Safety, NLP, Computational Social Science

HIGHLIGHTS

Funding: NSF CAREER Award (2025–2030, $700,000) to develop methods that help online communities reinforce norm-aligned behavior, not only punish harm (Siebel School News).

Recognition and Leadership: Best Paper, CSCW 2023; Best Paper nominations, EMNLP 2026 and PoliSim@CHI 2026; Editor's Pick, ACM TOCHI. Associate Editor, CSCW (2026–2027); Associate Chair, CHI (2024–2027). Keynote, Purdue University symposium (2024).

Teaching and Mentoring: Named to the *List of Teachers Ranked as Excellent by Their Students* four times; designed five new courses. Four PhD students have graduated from my SCUBA Lab (2025–2026), with placements including an Assistant Teaching Professor position at Purdue and positions at Google and Morgan Stanley; advisees have received the C.W. Gear Award and an NSF GRFP Honorable Mention.

Research and Policy Impact: A causal evaluation of Reddit's 2015 ban of hateful communities (CSCW 2017) found that users who stayed active on Reddit cut their hate speech by at least 80%. Reddit CEO Steve Huffman cited the study in written testimony to the U.S. House Energy and Commerce Committee (October 16, 2019); the findings informed Reddit's subsequent bans of hate groups and were included in its 2018 Transparency Report. Crossmod (CSCW 2019), an open-source moderation tool evaluated in a controlled study, was later deployed in a Reddit community of over 14 million subscribers.

HONORS & AWARDS

2026 Best Paper Nominee, EMNLP 2026 *(LLUMI)*
2026 Best Paper Nominee, PoliSim@CHI 2026 *(From Plausible to Causal)*
2025 NSF CAREER Award
2024 Editor's Pick for Notable Paper, ACM TOCHI *(Quarantined!)*
2024 Honorable Mention, DIS 2024 *(Computational Approaches to Addressing Online Harassment)*
2023 Best Paper, CSCW 2023 *(User-Moderator Alignment on r/ChangeMyView)*
2020 – 2022 List of Teachers Ranked as Excellent by Their Students, University of Illinois Urbana-Champaign *(Fall 2020, Spring 2021, Fall 2021, Spring 2022)*
2020 Best Paper Award Runner Up, WebSci 2020 *(Russian Troll Accounts on Twitter)*

EDUCATION

Aug 2015 – May 2020	**Ph.D. in Computer Science**, Georgia Institute of Technology, Atlanta, GA Thesis: *Combatting Abusive Behavior in Online Communities Using Cross-Community Learning* Committee: Eric Gilbert (advisor), Amy Bruckman, Munmun De Choudhury, Jacob Eisenstein, Cliff Lampe
2010 – 2015	**B.Tech & M.Tech in Computer Science and Engineering**, Indian Institute of Technology (IIT) Madras, Chennai, India Thesis: *Footprints on Silicon: Explorations in Gathering Autobiographical Content*

FUNDING

2025 – 2030 $700,000	PI. National Science Foundation (NSF) CAREER Award, *CAREER: Models and Methods to Help Online Communities Support Behaviors Aligned with Community Norms*.
2022 $5,000	PI. Title VI Award, European Union Center (EUC), University of Illinois, *Uses of Artificial Intelligence for Inclusive Online Interactions: New Course including European partners*.
2020 – 2021 $68,551	Illinois PI (lead PI: Eric Gilbert, University of Michigan). Facebook Research Award, Content Governance RFP, *P(Healthy) != 1-P(Toxicity): Measuring healthy online behavior*. Illinois share of a $100,000 total award.

PUBLICATIONS

*In computer science, top-tier conferences, in addition to journals, are venues for highly selective, double-blind, peer-reviewed, archival publications. Ph.D., master's, and undergraduate students supervised at Illinois are underlined; * denotes equal contribution.*

Peer-Reviewed Conference and Journal Papers

C.38 **Beyond Throughput: Objectives and Tradeoffs in Volunteer Moderation**
ICWSM 2027 (forthcoming)
Tanvi Bajpai, **Eshwar Chandrasekharan**

C.37 **VASTU: Language Models Struggle to Recognize Online Community Values**
EMNLP 2026 (Main)
Agam Goyal, Xianyang Zhan, Charlotte Lambert, Koustuv Saha, **Eshwar Chandrasekharan**

C.36 **LLUMI: Improving LLM Writing Assistance for Mental Health Support with Online Community Feedback**
EMNLP 2026 (Main) ***Nominated for Best Paper Award***
Jiwon Kim, Maya Ajit, Sherry Gong, Soorya Ram Shimgekar, Dong Whi Yoo, **Eshwar Chandrasekharan**, Koustuv Saha

C.35 **The Hidden Toll of Social Media News on Psychosocial Wellbeing**
EMNLP 2026 (Main)
Olivia Pal, Agam Goyal, **Eshwar Chandrasekharan**, Koustuv Saha

C.34 **Facts Without Rules: Boundary Metadata Collapse in Multi-Agent LLM Handoffs**
EMNLP 2026 (Findings)
Yian Wang, Agam Goyal, **Eshwar Chandrasekharan**, Hari Sundaram

C.33 **Answer Bubbles: Information Exposure in AI-Mediated Search**
EMNLP 2026 (Findings)
Michelle Huang*, Agam Goyal*, Koustuv Saha, **Eshwar Chandrasekharan**

C.32 **The Language of Approval: Identifying the Drivers of Positive Feedback Online**
CHI 2026
Agam Goyal, Charlotte Lambert, **Eshwar Chandrasekharan**

C.31 **"Think about it like you're a firefighter": Understanding How Reddit Moderators Use the Modqueue**
CHI 2026
Tanvi Bajpai, **Eshwar Chandrasekharan**

C.30 **Needling Through the Threads: A Visualization Tool for Navigating Threaded Online Discussions**
CHI 2026
Yijun Liu, Frederick Choi, **Eshwar Chandrasekharan**

C.29 **Masking or Mitigating? Deconstructing the Impact of Query Rewriting on Retriever Biases in RAG**
ACL 2026
Agam Goyal, Koyel Mukherjee, Apoorv Saxena, Anirudh Phukan, **Eshwar Chandrasekharan**, Hari Sundaram

C.28 **Examining Algorithmic Curation on Social Media: An Empirical Audit of Reddit's r/popular Feed**
ICWSM 2026
Jackie Chan, Frederick Choi, Koustuv Saha, **Eshwar Chandrasekharan**

C.27 **Uncovering the Internet's Hidden Values: An Empirical Study of Desirable Behavior Using Highly-Upvoted Content on Reddit**
ICWSM 2026
Agam Goyal, Charlotte Lambert, Yoshee Jain, **Eshwar Chandrasekharan**

C.26 **MoMoE: Mixture of Moderation Experts Framework for AI-Assisted Online Governance**
EMNLP 2025 (Main)
Agam Goyal, Xianyang Zhan, Yilun Chen, Koustuv Saha, **Eshwar Chandrasekharan**

C.25 **ArgCMV: An Argument Summarization Benchmark for the LLM-era**
EMNLP 2025 (Main)
Omkar Gurjar, Agam Goyal, **Eshwar Chandrasekharan**

C.24 **Venire: A Machine Learning-Guided Panel Review System for Community Content Moderation**
CSCW 2025
Vinay Koshy, Frederick Choi, Yi-Shyuan Chiang, Hari Sundaram, **Eshwar Chandrasekharan**, Karrie Karahalios

C.23 **The Chilling: Identifying Strategic Antisocial Behavior Online and Examining the Impact on Journalists**
CSCW 2025
Yian Wang, Mukhilshankar Umashankar, **Eshwar Chandrasekharan**, Hari Sundaram

C.22 **Does Positive Reinforcement Work?: A Quasi-Experimental Study of the Effects of Positive Feedback on Reddit**
CHI 2025
Charlotte Lambert, Koustuv Saha, **Eshwar Chandrasekharan**

C.21 **Creator Hearts: Investigating the Impact of Positive Signals from YouTube Creators in Shaping Comment Section Behavior**
CHI 2025
Frederick Choi, Charlotte Lambert, Vinay Koshy, Sowmya Pratipati, Tue Do, **Eshwar Chandrasekharan**

C.20 **SLM-Mod: Small Language Models Surpass LLMs at Content Moderation**
NAACL 2025 (Main)
Xianyang Zhan*, Agam Goyal*, Yilun Chen, **Eshwar Chandrasekharan**, Koustuv Saha

C.19 **Measuring epistemic trust: Towards a new lens for democratic legitimacy, misinformation, and echo chambers**
CSCW 2024
Dominic Zaun Eu Jones, **Eshwar Chandrasekharan**

C.18 **"Positive reinforcement helps breed positive behavior": Moderator Perspectives on Encouraging Desirable Behavior**
CSCW 2024
Charlotte Lambert, Frederick Choi, **Eshwar Chandrasekharan**

C.17 **Understanding Community Resilience: Quantifying the Effects of Sudden Popularity via Algorithmic Curation**
ICWSM 2024
Jackie Chan, Charlotte Lambert, Frederick Choi, Stevie Chancellor, **Eshwar Chandrasekharan**

C.16 **Opportunities, Tensions, and Challenges in Computational Approaches to Addressing Online Harassment**
DIS 2024 ***Honorable Mention***
Evey Jiaxin Huang, Abhraneel Sarma, Sohyeon Hwang, **Eshwar Chandrasekharan**, Stevie Chancellor

C.15 **ConvEx: A Visual Conversation Exploration System for Discord Moderators**
CSCW 2023
Frederick Choi, Tanvi Bajpai, Sowmya Pratipati, **Eshwar Chandrasekharan**

C.14 **Measuring User-Moderator Alignment on r/ChangeMyView**
CSCW 2023 ***Best Paper***
Vinay Koshy, Tanvi Bajpai, **Eshwar Chandrasekharan**, Hari Sundaram, Karrie Karahalios

C.13 **Conversational Resilience: Quantifying and Predicting Conversational Outcomes Following Adverse Events**
ICWSM 2022
Charlotte Lambert, Ananya Rajagopal, **Eshwar Chandrasekharan**

C.12 **Harmonizing the Cacophony with MIC: An Affordance-aware Framework for Platform Moderation**
CSCW 2022
Tanvi Bajpai, Drshika Asher, Anwesa Goswami, **Eshwar Chandrasekharan**

C.11 **Quarantined! Examining the Effects of a Community-Wide Moderation Intervention on Reddit**
TOCHI 2022 ***Editor's Pick for Notable Paper, 2024***
Eshwar Chandrasekharan, Shagun Jhaver, Amy Bruckman, Eric Gilbert

C.10 **Conversations Gone Alright: Quantifying and Predicting Prosocial Outcomes in Online Conversations**
WWW 2021
Jiajun Bao, Junjie Wu, Yiming Zhang, **Eshwar Chandrasekharan** and David Jurgens

C.9 **Synthesized Social Signals: Computationally-Derived Social Signals from Account Histories**
CHI 2020
Jane Im, Sonali Tandon, **Eshwar Chandrasekharan**, Taylor Denby, Eric Gilbert

C.8 **Still out there: Modeling and Identifying Russian Troll Accounts on Twitter**
WebSci 2020 ***Best Paper Award Runner Up***
Jane Im, **Eshwar Chandrasekharan**, Jackson Sargent, Paige Lighthammer, Taylor Denby, Ankit Bhargava, Libby Hemphill, David Jurgens, Eric Gilbert

C.7 **Crossmod: A Cross-Community Learning-based System to Assist Reddit Moderators**
CSCW 2019
Eshwar Chandrasekharan, Chaitrali Gandhi, Matthew Wortley Mustelier, Eric Gilbert

C.6 **A Just and Comprehensive Strategy for Using NLP to Address Online Abuse**
ACL 2019
David Jurgens, **Eshwar Chandrasekharan**, Libby Hemphill

C.5 **Prevalence and Psychological Effects of Hateful Speech in Online College Communities**
WebSci 2019
Koustuv Saha, **Eshwar Chandrasekharan**, Munmun De Choudhury

C.4 **The Internet's Hidden Rules: An Empirical Study of Reddit Norm Violations at Micro, Meso, and Macro Scales**
CSCW 2018
Eshwar Chandrasekharan, Mattia Samory, Shagun Jhaver, Hunter Charvat, Amy Bruckman, Cliff Lampe, Jacob Eisenstein, Eric Gilbert

C.3 **You Can't Stay Here: The Efficacy of Reddit's 2015 Ban Examined Through Hate Speech**
CSCW 2017
Eshwar Chandrasekharan, Umashanthi Pavalanathan, Anirudh Srinivasan, Adam Glynn, Jacob Eisenstein, Eric Gilbert

C.2 **Situated Anonymity: Impacts of Anonymity, Ephemerality, and Hyper-Locality on Social Media**
CHI 2017
Ari Schlesinger, **Eshwar Chandrasekharan**, Christina Masden, Amy Bruckman, W Keith Edwards, Rebecca Grinter

C.1 **The Bag of Communities: Identifying Abusive Behavior Online with Preexisting Internet Data**
CHI 2017
Eshwar Chandrasekharan, Mattia Samory, Anirudh Srinivasan, Eric Gilbert

Preprints Under Review

P.13 **First Impressions: How Placement Shapes the Influence of AI Summaries**
(under review)
Claire Wang, Agam Goyal, Frederick Choi, Koustuv Saha, **Eshwar Chandrasekharan**

P.12 **The Wisdom of the Loudest: A Large-Scale Audit of Generative Search on Reddit**
(under review)
Agam Goyal, Claire Wang, **Eshwar Chandrasekharan**

P.11 **Stereotypically Yours: Portrayal and Perception of Race-Coded AI Companions**
(under review)
Claire Wang, Jiayue Melissa Shi, Agam Goyal, Grace E Sletten, Renwen Zhang, **Eshwar Chandrasekharan**, Koustuv Saha

P.10 **PeerPen: AI-Assisted Writing for Online Mental Health Peer Support**
(under review)
Jiwon Kim, Sherry Gong, Maya V Ajit, Soorya Ram Shimgekar, Yunhao Yuan, Dong Whi Yoo, **Eshwar Chandrasekharan**, and Koustuv Saha

P.9 **Trustworthy FinAInce: Unpacking How AI-Mediated Financial Advice is Judged**
(under review)
Aryan Ramchandra Kapadia, **Eshwar Chandrasekharan**, and Koustuv Saha

P.8 **X-pressing Intimacy: Mental Health Disclosure After Twitter Acquisition**
(under review)
Olivia Pal, Yunhao Yuan, Meghana Mandava, Talayeh Aledavood, **Eshwar Chandrasekharan**, Koustuv Saha

P.7 **Prompt PerturbAItions: Mechanisms of LLM Response Inconsistencies**
(under review)
Soorya Ram Shimgekar, Agam Goyal, Amruta Parulekar, Joshua Chen, Yian Wang, Navin Kumar, Hari Sundaram, **Eshwar Chandrasekharan**, Koustuv Saha

P.6 **Algorithmic Cultivation: How Social Media Feeds Shape User Language**
(under review)
Olivia Pal, Agam Goyal, **Eshwar Chandrasekharan**, Koustuv Saha

P.5 **Social Simulacra in the Wild: AI Agent Communities on Moltbook**
(under review)
Agam Goyal, Olivia Pal, Hari Sundaram, **Eshwar Chandrasekharan**, Koustuv Saha

P.4 **PAIR-SAFE: A Paired-Agent Approach for Runtime Auditing and Refining AI-Mediated Mental Health Support**
(under review)
Jiwon Kim, Violeta J. Rodriguez, Dong Whi Yoo, **Eshwar Chandrasekharan**, Koustuv Saha

P.3 **The Ranking Effect: How Position Shapes Attention on Trending Feeds**
(under review)
Jackie Chan, Frederick Choi, Koustuv Saha, **Eshwar Chandrasekharan**

P.2 **PILOT: Control Surfaces for Authoring Social Media Feeds**
(under review)
Agam Goyal, Frederick Choi, **Eshwar Chandrasekharan**

P.1 **Mind Your Ps and Qs: Positive Moderation Practice in the Positive Queue**
(under review)
Charlotte Lambert, Agam Goyal, Eunice Mok, **Eshwar Chandrasekharan**

Posters, Extended Abstracts, Workshop Papers, and Other Publications

L.11 **From Plausible to Causal: Counterfactual Semantics for Policy Evaluation in Simulated Online Communities**
PoliSim@CHI 2026 Best Paper Nominee
Agam Goyal, Yian Wang, **Eshwar Chandrasekharan**, Hari Sundaram

L.10 **Voices Lost in the Algorithm: Improving Instagram's Moderation of Mental Health Communities**
CSCW Companion 2025 (poster)
Eunice Mok*, Yoshee Jain*, Sriya Gottiparthi*, **Eshwar Chandrasekharan**

L.9 **Investigating How Gilds Were Employed on Reddit**
CSCW Companion 2024 (poster)
Charlotte Lambert, Yoshee Jain, Koustuv Saha, **Eshwar Chandrasekharan**

L.8 **LLM-Mod: Can Large Language Models Assist Content Moderation?**
CHI Late Breaking Work 2024
Mahi Kolla, Siddharth Salunkhe, **Eshwar Chandrasekharan**, Koustuv Saha

L.7 **What's in a Social Computing Course: Analyzing Computer and Information Science Syllabi**
EduCHI '24: Proceedings of the 6th Annual Symposium on HCI Education
Catherine Delcourt, Sukrit Venkatagiri, **Eshwar Chandrasekharan**

L.6 **The Future of Research on Social Technologies: CCC Workshop Visioning Report**
Computing Community Consortium Workshop 2023, The Future of Research on Social Technologies
Motahhare Eslami, Eric Gilbert, Sarita Schoenebeck, Eric P. S. Baumer, **Eshwar Chandrasekharan**, Michelle De Mooy, Karrie Karahalios, David Karger, Tressie McMillan Cottom, Andrés Monroy-Hernández, Loren Terveen, John Wihbey

L.5 **Future of Workplace Conversational AI-driven Personal Assistants: Promises and Perils**
FIRE '23: Proceedings of the 15th Annual Meeting of the Forum for Information Retrieval Evaluation
Koustuv Saha, **Eshwar Chandrasekharan**

L.4 **Evolution of Rules in Reddit Communities**
CSCW Companion 2023 (poster)
Harita Reddy, **Eshwar Chandrasekharan**

L.3 **Community Resilience: Quantifying the Disruptive Effects of Sudden Spikes in Activity within Online Communities**
CHI Late Breaking Work 2022.
Jackie Chan, Aditi Atreyasa, Stevie Chancellor, **Eshwar Chandrasekharan**

L.2 **Hybrid Approaches to Detect Comments Violating Macro Norms on Reddit**
arXiv. Dataset paper 2019.
Eshwar Chandrasekharan, Eric Gilbert

L.1 **Footprints on Silicon: Explorations in Gathering Autobiographical Content**
CICLing 2015 (IJCLA)
Eshwar Chandrasekharan, Sutanu Chakraborti

Special Interest Groups Organized

S.1 **Community-Driven Models for Research on Social Platforms**
CSCW Companion '24 *[SIG at CSCW 2024]*
Sanjay Kairam, Michael S. Bernstein, Amy Bruckman, Stevie Chancellor, **Eshwar Chandrasekharan**, Munmun De Choudhury, Casey Fiesler, Hanlin Li, Nicholas Proferes, Manoel Horta Ribeiro, C. Estelle Smith, Galen Weld

Workshops Organized

W.1 **Volunteer Work: Mapping the Future of Moderation Research**
CSCW 2019*. [Accepted as a* ***full-day workshop****].* ***ACM Extended Abstract.***
Charles Kiene, Shagun Jhaver, **Eshwar Chandrasekharan**, Bryan Dosono, Brianna Dym, Sarah Gilbert, Jialun "Aaron" Jiang, Kat Lo, Joseph Seering, Kenny Shores, Donghee "Yvette" Wohn

W.2 **Managing and Designing for Norms in Online Communities**
ICWSM 2019*. [Accepted as a* ***full-day workshop****]*
Joseph Seering, Sanjay Kairam, **Eshwar Chandrasekharan**, Stevie Chancellor

SELECTED PRESS

2025 **Siebel School News**, *CS professor Eshwar Chandrasekharan receives NSF CAREER Award*.

2025 **The New York Times**, *'Will I Lose My Job?' Federal Workers Flock to Reddit for Answers*.

2019 **The Verge**, *Why deplatforming 8chan can be effective*.

2019 **MIT Technology Review**, *CloudFlare dropping 8chan helps fight hate even if 8chan comes back*.

2017 **The New York Times**, *Reddit Limits Noxious Content by Giving Trolls Fewer Places to Gather*.

2017 **TechCrunch**, *Study finds Reddit's controversial ban of its most toxic subreddits actually worked*.

2017 **Georgia Tech News Center**, *New Georgia Tech Research May Help Combat Abusive Online Comments*.

INDUSTRY EXPERIENCE

May – Aug 2018 **Health ML Intern, Engineering**, Twitter Inc.

- Examined the landscape of incoming abuse reports on Twitter in a large-scale, empirical manner to quantify the notion of *report quality*.
- Identified the different types of abuse reported by users using a mixed-methods approach (computational and qualitative).
- Proposed approaches for building new ML tools to detect and discourage different types of abuse on the platform, and for training existing internal tools in a more nuanced manner.
- Area of work: Machine Learning, Quantitative Research, Data Science.

May – Aug 2017 **Data Science Intern**, BuzzFeed Inc.

- Developed a human-machine collaboration service that informs BuzzFeed's content distribution strategy on Facebook.
- Built machine learning models to predict which BuzzFeed links would perform best on certain Facebook pages, and to identify which pages to post on for the most content shares.
- Area of work: Data Science, Machine Learning.

May – Jul 2013 **Software Development Engineer (SDE) Intern**, Amazon Development Center (India) Pvt. Ltd., Bangalore

- Automated product opinion summarization of Amazon customer reviews.
- Created original content for product web pages to improve search engine optimization.
- Area of work: Opinion Mining, Sentiment Analysis, Text Summarization.

TEACHING

University of Illinois Urbana-Champaign, Course Instructor

Fall 2026	CS598 CSS: Computational Social Science
Fall 2025	CS598 OM: Online Moderation *(newly designed)*
Spring 2025	CS498 SC: Computational Social Science
Fall 2024	CS598 CSS: Computational Social Science
Spring 2024	CS498 SC: Computational Social Science *(newly designed)*
Fall 2023	CS598 SPH: Research Methods in HCI *(newly designed)*
Fall 2022	CS498 ASU: Antisocial Computing
Spring 2022	CS598 SC: Social Computing
Fall 2021	CS598 ECH: Antisocial Computing
Spring 2021	CS598 EWS: Social Computing *(newly designed)*
Fall 2020	CS598 RBO: Antisocial Computing *(newly designed)*

Georgia Institute of Technology

Jan – Dec 2017 Minor in Higher Education, Center for Teaching and Learning (CETL)

INVITED TALKS, PANELS & GUEST LECTURES

Oct 2026	**Panelist, COMPASS Workshop Discussion Panel** Socio-Technical Approaches to Content Moderation & Platform Governance, CSCW 2026
Sep 2024	**"Proactive Approaches to Promote Community Resilience and Foster Desirable Behavior Online"** Keynote Address, Symposium on Frontiers in Online Community Research, Purdue University
Fall 2024	**"Advances in Content Moderation Research"** Guest Lecture, Social Spaces on the Internet (CS 567), University of Illinois Urbana-Champaign
Jun 2023	**"Proactive Approaches to Online Governance"** Invited Talk, HCIC Annual Meeting on Supporting Online Trust and Safety, Privacy, and Security
Mar 2023	**Panelist, "New Work on Community Driven Governance"** Social Media Governance Initiative Spring 2023 Convening, Yale Law School
Feb 2022	**"Quantifying Outcomes in Online Conversations"** Guest Speaker, Jigsaw, Google
Sep 2021	**Panelist, "Transparent Automated Content Moderation"** TU Vienna Digital Humanism and Illinois EU Center Faculty Workshop, University of Illinois
Feb 2021	**"Understanding, building and evaluating social computing systems"** Guest Lecture, HCI Methods, University of Minnesota
Oct 2020	**Panelist, "Stories of refusal from within the corporate firm"** Panel, The Refusal Conference, AFOG, UC Berkeley
Sep 2020	**"Combating trolling and online harassment"** Guest Lecture, Social Spaces on the Internet (CS598kgk), University of Illinois

Jul 2020	**"Quantifying Prosocial Outcomes in Online Conversations"** Invited Talk, Stanford HCI Group, Stanford University
Feb – Apr 2020	**"Using Human-Centered AI to Combat Abusive Behavior on the Internet"** Invited talks: UMass Amherst, College of Information and Computer Sciences (Apr 2020); University of Illinois, Department of Computer Science (Mar 2020); University of Minnesota, Department of Computer Science (Feb 2020); UC Berkeley, Department of EECS (Feb 2020)
Feb 2019	**"Combating Online Abusive Behavior Using Cross-Community Learning"** Talk, School of Information, University of Michigan, Ann Arbor, MI
Jun 2018	**"Effectiveness of *Banning* as a Moderation Strategy"** Talk, Twitter HQ, San Francisco, CA
Apr 2015	**"Twitter Data Mining and Analysis"** Guest Lecture, Social Computing (CS6474), Georgia Institute of Technology
Jul 2014	**"Mining the Web to predict Future Events"** Guest Lecture, Artificial Intelligence & Database Laboratory, IIT Madras

SERVICE

Editorial Service and Senior Program Committee

2026 – 2027	Associate Editor, ACM Conference on Computer-Supported Cooperative Work and Social Computing (CSCW)
2024 – 2027	Associate Chair, ACM CHI Conference on Human Factors in Computing Systems (CHI)
2020 – 2021	Associate Chair, ACM Conference on Computer-Supported Cooperative Work and Social Computing (CSCW)

Conference Organization

2024	Special Interest Group (SIG) co-organizer, CSCW 2024 (see S.1)
2022	Data Challenge Chair, ICWSM 2022
2019	Workshop co-organizer, CSCW 2019 and ICWSM 2019 (see W.1, W.2)
2019	Local Arrangements Co-Chair, CSCW 2019

Program Committee and Reviewing

2018 – 2020	Program Committee: EMNLP 2020, WWW 2020, ALW 2019 (ACL'19), TRAC-1 (COLING'18)
2017 – 2022	Reviewer: CHI 2017–'22, CSCW 2017–'20 & 2021–'22, ICWSM 2018–'20, ACL 2019–'20, TOCHI 2018, TSC 2019–'20

Grant Review

2026	Panelist, NSF CISE Proposal Review Panel, Human-Centered Computing (HCC)

Departmental Service, University of Illinois Urbana-Champaign

2025 – 26	PhD Launch Committee
2024 – 25	Graduate Admissions Committee
2023 – 24	Graduate Admissions Committee
2022 – 23	Tenure-Track Recruiting Committee; Academic Appeals Committee
2021 – 22	Tenure-Track Recruiting Committee
2020 – 21	FAA (Graduate Admissions)

Doctoral Committees

Prelim and defense	Yian Wang: Prelim (2026) Jackie Chan: Prelim (2024–2025), Defense (2026) (advisor) Frederick Choi: Prelim (2024), Defense (2026) (advisor) Charlotte Lambert: Prelim (2024), Defense (2026) (advisor) Tanvi Bajpai: Prelim (2024), Defense (2025) (advisor) Vinay Koshy: Prelim (2023), Defense (2025) Hyo Jin Gina Do: Prelim (2021), Defense (2022)
Qualifying exam	Mung Yao Jia (2026), Indu Panigrahi (2026), Yifan Song (2025), Ziheng Huang (2025), Bhavana Bheem (2025), Alex Atcheson (2023), Ali Zaidi (2022), Muhammad Rizky Wellyanto (2022), Pingjing Yang (2021), Ti-Chung Cheng (2021), Vinay Koshy (2021), Charlotte Yoder (2021)

MENTORING

† = Awarded CS STARS at Illinois. *Next position* = first role after working with me; *Current position* = role as of 2026.

Current Students at Illinois

Fall 2024 –	Agam Goyal (PhD)
Fall 2025 –	Claire Wang (PhD)
2025 –	Eunice Mok† (Undergraduate). Honorable Mention, 2025–26 CRA Outstanding Undergraduate Researcher Award

PhD Alumni

2021 – 2026	**Charlotte Lambert**. *Next position:* Assistant Teaching Professor, Purdue University *Honors:* C.W. Gear Outstanding Graduate Student Award (2023); Outstanding TA Lifetime Award (2024); Saburo Muroga Endowed Fellowship (2020–21); Ranked as Excellent by Students (3 semesters) Thesis: *Fostering Healthier Online Communities through Positive Approaches to Moderation*
2021 – 2026	**Jackie Chan**. *Next position:* Assistant Vice President, Morgan Stanley *Honors:* NSF Graduate Research Fellowship Program (GRFP) Honorable Mention Thesis: *Quantifying the Impacts of Algorithmic Ranking on Social Media Trending Feeds*
2021 – 2026	**Frederick Choi** *Honors:* Roy J. Carver Fellowship (2020–21) Thesis: *Designing Technologies to Support Moderation and Curation of Online Spaces at the Community and Individual Level*
2021 – 2025	**Tanvi Bajpai**. *Next position:* Software Engineer, Google *Honors:* Outstanding TA Award, Illinois CS (Fall 2020) Thesis: *Beyond the Black Box: Towards Bridging Theoretical Computer Science and Human-Computer Interaction*

MS Alumni

2024 – 2025 **Omkar Gurjar**. *Current position:* Member of Technical Staff, Microsoft AI
Thesis: *Distilling Arguments: A Study of Human and LLM Persuasion in the Online Discourse*

2022 – 2023 **Dominic Jones**. *Next position:* Data Scientist, Reserve Bank of Australia
Thesis: *Measuring Epistemic Trust: Towards a New Lens for Democratic Legitimacy, Misinformation, and Echo Chambers*

2022 – 2023 **Harita Reddy**
Thesis: *Studying Rule Trajectories in Reddit Communities*

Undergraduate Alumni

† = Awarded CS STARS at Illinois

- Yoshee Jain†. *Next position:* PhD student, University of Illinois Urbana-Champaign
- Sowmya Pratipati†. *Next position:* Software Engineer, Stripe
- Xianyang Zhan. *Next position:* Software Engineer, Microsoft
- Aarohi Gupta. *Next position:* M.S. in Computer Science, Stanford University
- Drshika Asher†. *Next position:* Research Software Engineer, National Center for Supercomputing Applications (NCSA)
- Ayaan Shah. *Next position:* Member of Technical Staff (ML), Prophet Security
- Tue Do. *Next position:* PhD student, University of Illinois Urbana-Champaign
- Ishq Gupta. *Next position:* Data Analyst, Capital One
- Ananya Rajagopal. *Next position:* Software Engineer, Chime
- Anwesa Goswami. *Next position:* AI/ML Infrastructure Engineer, JPMorganChase
- Jacqueline Chen. *Next position:* Software Engineer, Gap Inc.
- Also: Yilun (Allen) Chen, Aarthi Balaji, Clark Lu, Niharika Chintalapati, Sriya Gottiparthi†

Students Mentored Before Illinois (as a PhD student)

- Anirudh Srinivasan (Georgia Tech, MS-CS). *Current position:* Director of AI, Ittiam
- Sonali Tandon (Michigan School of Information, MS). *Current position:* Senior Product Designer, Roblox
- Chaitrali Gandhi (Michigan School of Information, MS). *Current position:* AI Product Management, ServiceNow
- Aditya Manjusha Shylesh (Michigan School of Information, BS). *Current position:* Software Engineer, Yahoo
- *Georgia Tech BS-CS*: Nancy Tao, Christine Shih, Iris Birman. *Michigan MS-SI*: Aljoharah Alfayez. *Michigan BS-SI*: Hunter Charvat, Katherine Mustelier, Yiming Zhang, Jiajun Bao, Junjie Wu, Michelle Lee, Shradha Reddy, Weiji Li, Robert Kang